

Adaptive and Power-Aware Resilience for Extreme-scale Computing

Anonymized

ABSTRACT

With the concerted efforts from researchers in hardware, software, algorithm, and data management, HPC is moving towards its next generation—extreme-scale computing. Featuring a computing capacity of quintillion (10^{18}) FLOPS, extreme-scale computing will be a significant achievement in computer science and engineering. As we approach the new era of computing, however, several daunting scalability challenges need to be conquered. Delivering extreme-scale performance will require a billion-core computing infrastructure, along with a dramatic increase in the number of memory modules, communications devices and storage components. At such a large scale, failure would become a norm instead of exception, driving the system to significantly lower efficiency with unprecedented amount of power consumption.

To tackle this challenge, we propose a proactive, power-aware resilience framework, referred to as Lazy Shadowing, as an efficient and scalable approach to achieve high-levels of resilience, through forward progress, in extreme-scale, failure-prone computing environments. Lazy Shadowing associates with each process a “shadow” process that executes at a reduced rate, and overlaps the recovery time after each failure with the time needed to roll forward the shadows to a consistent state. Compared to existing fault tolerance methods, our approach can achieve 20% energy saving with reduced solution time at scale.

Categories and Subject Descriptors

C.4 [Performance of systems]: Fault tolerance; C.2.4 [Distributed Systems]: Distributed applications

General Terms

Design, Performance, Reliability

Keywords

Lazy Shadowing, extremely-scale computing, forward progress

1. INTRODUCTION

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

SC '15 Austin, TX, USA

Copyright 20XX ACM X-XXXXX-XX-X/XX/XX ...\$15.00.

To enable future scientific breakthroughs and discoveries, the next generation of scientific applications will require extreme-scale computing performance to support the execution of predictive models and analysis of massive quantities of data, with significantly higher resolution and fidelity than what can be achieved in existing computing infrastructure. In order to achieve the desired functional performance of these applications, current HPC supercomputers need to increase the number of cores by orders of magnitude while also requiring a dramatic increase in the memory modules, communications devices and storage components [2].

The explosive growth in the number of components significantly exacerbates the propensity of extreme-scale computing systems to faults. Regardless of the individual components' reliability, system resilience will continue to decrease as the number of components increases, as shown in Figure 1. For example, a system of 1 million components, each averaging a fault every 25 years, is expected to experience a fault every 10 minutes.

To avoid the full re-execution of an application upon failure, traditional fault tolerance methods periodically checkpoint the execution state of the applications, such that recovery can be achieved by restarting from a consistent checkpoint when failure occurs [19]. Recent works [22, 28] have shown that such solutions are unlikely to scale to the level of faults anticipated in extreme-scale environments. Given the increase in system failure rates and the time required to checkpoint large-scale compute- and data-intensive applications, it is very likely that the time spent on checkpointing and recovery may exceed the system mean time between failures, thereby reducing considerably the overall performance of the system. Furthermore, the nature and frequency of errors in extreme-scale computing are such that checkpoint-restart may not be an adequate approach to handle the diverse types of faults in these environments. Several studies have shown the same concern and proposed process replication as a possible solution [14, 28].

Process replication, also referred to as state-machine replication, is a well-known technique that has been shown to scale to meet the resilience requirements of large distributed and mission-critical systems [29]. Based on this technique, a process's state and computation are replicated across independent computing nodes. When the main process fails, one of the replicas takes over the computation task. Replication requires at least doubling the number of computing resources, since each process must have at least 1 replica, thereby limiting the system efficiency to 50%.

Another direct consequence of the increase in system scale is the dramatic increase in power requirements. Analysis in [2] shows a steady rise of system power consumption to 1-3MW in 2008, followed by a sharp increase to 10-20MW in subsequent years, with the expectation that it could surpass 50MW by 2016, making system power a leading design constraint on the path to extreme-scale




Figure 1: Impact of system size on system resilience.

computing. The US Department of Energy has recognized this trend and established a power limit of 20 megawatt [2], challenging the research community to provide a 1000x improvement in performance with only a 10x increase in power. Although largely pragmatic¹, adhering to this power constraint is critical to the design and operations of extreme-scale computing systems. The power-limits placed upon extreme-scale systems is at the heart of the tradeoff between power consumption and time-to-solution. Although the intricacies of this tradeoff are complex, operating under power-constraints is a reality that future extreme-scale system designers must face, making the relationship between computing performance and power consumption increasingly more critical.

Insofar as performance is concerned, resilience to failures and adherence to power constraints are two conflicting objectives, as achieving high performance may push the system’s components past their thermal limit and increase their likelihood of failure. The inherent instability of extreme-scale computing systems, in terms of the envisioned high-rate and diversity of their faults, together with the demanding power constraints under which these systems will be designed to operate, calls for a reconsideration of the fault tolerance problem. The main objective of this paper is to explore novel paradigms for resilience in future extreme-scale systems.

To this end, we propose a proactive, power-aware resilience model, referred to as *Lazy Shadowing*, as an efficient and scalable alternative to achieve high-levels of resilience, through forward progress, in extreme-scale, failure-prone computing environments. In the proposed resilience model, each process is associated with a replica, which runs concurrently with its main process. For clarity, the original process is referred to as the main (process) and the replica as the shadow (process). The shadow executes the same code as its associated main process, but at a reduced CPU rate and on a different computing node. The rate at which the shadow runs is derived based on the expected time-to-solution of the supported application, the power constraints imposed by the underlying computing infrastructure, and the likelihood of failure. Upon failure of the main process, the shadow increases its execution rate to complete the task, thereby reducing the impact of such a failure on the progress of the remaining tasks. Successful completion of the main process, however, results in immediate termination of the shadow. Since the failure rate of an individual component is much lower than that of the whole system, it is very likely that, most of the main processes complete their execution successfully. Consequently, the coupled execution of a main process and its associated shadow dra-

matically increases a power-constrained system’s tolerance to failure, at a significantly reduced energy consumption.

The rest of the paper is organized as follows. We begin with a survey on related work in Section 2. Section 3 defines the Lazy Shadowing computational model and Section 4 discusses how Lazy Shadowing can be used in extreme-scale computing environments for efficient fault tolerance. We then introduce three analytical models for performance assessment in Section 5, followed by experiments and evaluation in section 6. Section 7 concludes this work and points out future research directions.

2. RELATED WORK

Rollback and recovery are predominate mechanisms to achieve fault tolerance in current HPC environments. In the most general form, the rollback and recovery mechanism involves the periodic saving of the current system state to stable storage, with the anticipation that in the case of a system failure, computation can be restarted from the most recently saved state [11]. Coordinated Checkpoint is a method to achieve a globally consistent state. Specifically, all processes coordinate with one another to produce individual states that satisfy the “happens before” communication relationship [6], which is proved to provide a consistent global state. The major benefit of coordinated checkpointing stems from its simplicity and ease of implementation. Its major drawback, however, is the lack of scalability, as it requires global process coordination to achieve a checkpoint [12, 28, 17].

In uncoordinated checkpointing, processes record their states independently and postpone creating a globally consistent view until the recovery phase [26]. The major advantage is the reduced overhead during fault free operation. However, the scheme requires that each process maintains multiple checkpoints and message logs, necessary to construct a consistent state during recovery. It can also suffer the well-known domino effect, which may result in the re-execution of the entire application [27]. One hybrid approach, known as communication induced checkpointing schemes, aims at reducing coordination overhead [3]. The approach, however, may cause processes to store useless states that are never used in future rollbacks. To address this shortcoming, “forced checkpoints” have been proposed [18]. This approach may lead to unpredictable checkpointing rates. Although well-explored, uncoordinated checkpointing has not been widely adopted in HPC environments, due to its dependency on applications [31, 15].

One of the largest overheads in any checkpointing process is the time necessary to write the checkpointing to stable storage, during which the application pause its execution. Incremental checkpointing attempts to address this shortcoming by only writing the changes since the previous checkpoint [7, 1]. This can be achieved using dirty-bit page flags [25, 13]. Hash based incremental checkpointing, on the hand, makes use of hashes to detect changes [7, 1]. Another proposed scheme, known as copy-on-write, offloads the checkpointing process to a secondary task and only writes incremental checkpoints [21]. The main concern of these techniques is that some applications would require increase in their memory to support the simultaneous execution of the checkpointing and the application. It has been suggested that nodes in extreme-scale systems should be configured with fast local storage, which improves the performance of checkpointing [2]. Multi-level checkpointing, which consists of writing checkpoints to multiple storage targets, can benefit from such a strategy [23, 16]. This, however, may lead to increased failure rates of individual nodes and increased per-node cost [8]. Furthermore, it may complicate the checkpoint writing process and requires that the system track the current location of all process’s checkpoints.

¹The constraint is derived from an estimated cost of \$1 million per MW per year and a power budget of \$20 million per year.

Process and state machine replication have long been used to provide fault tolerance in distributed and mission critical systems. To ensure consistent system state, all process messages must be delivered to all nodes running a given process, typically done by using a message ordering protocol [20]. Replication can be used to detect and correct system failures that are otherwise undetectable, such as silent data failures [24] and Byzantine faults [14]. Full and partial replication have also been proposed, to augment existing checkpointing techniques, and to guard against silent data corruption [30, 10]. To the best of our knowledge, Lazy Shadowing is the first attempt to explore a state-machine replication based framework that provides a tradeoff between time and hardware redundancy while meeting resilience and power requirements.

3. LAZY SHADOWING

We use the term core to represent the computing resource allocation unit [5]. We further use $P(\sigma, w, t)$ to denote a process executing at rate σ to complete a workload w by time t . The basic tenet of the Lazy Shadowing resilience model is the concept of shadowing, whereby each process is associated with a *lazy* replica. Assuming a single failure, Lazy Shadowing can be described as follows:

- A main process, $P_m(\sigma_m, w, t_m)$, that executes at the rate of σ_m to complete a task of size w at time t_m .
- A shadow process, $P_s(< \sigma_s^b, \sigma_s^a >, w, t_s)$, that executes at σ_s^b before failure, and σ_s^a after failure, to complete a task of size w at time t_s .


The main process and its shadow start executing simultaneously on two different cores. The main process executes at rate σ_m , while the shadow initially executes at rate $\sigma_s^b \leq \sigma_m$. In the absence of failure, the main process completes execution at time $t_m = w/\sigma_m$, which immediately triggers the termination of the shadow. However, if at time $t_f < t_m$ the main process fails, the shadow, which has completed an amount of work $w_b = \sigma_s^b * t_f$, increases its execution rate to σ_s^a to complete the task by t_s . The execution dynamics of Lazy Shadowing are depicted in Figure 2.

The execution rates of the shadow, σ_s^b and σ_s^a , can be derived by balancing the trade-offs between completion time and energy consumption. For a delay-tolerant, energy-stringent application, σ_s^b is set to 0, and the shadow starts executing only upon failure of the main process. For a delay-stringent, energy-tolerant application, the shadow starts executing at $\sigma_s^b = \sigma_m$ to guarantee the completion of the task at the specified time t_m , regardless of when the failure occurs. For delay and energy tolerant applications, a broad spectrum of delay and energy tradeoffs can be explored either empirically or using optimization frameworks.


To control the shadow's execution rate, Dynamic Voltage and Frequency Scaling (DVFS), a commonly used power management technique, can be applied. The effectiveness of DVFS, however, may be markedly reduced in computational platforms that exhibit saturation of the processor clock frequencies, large static power consumption, or small power dynamic range. An alternative to DVFS is to collocate multiple processes on a single core, while keeping the core running at maximum rate. Time sharing can then be used to achieve the desired execution rates of the processes. Given our focus on extreme-scale, multi-core computing infrastructure, we will use collocation for execution rate control.

4. EXTREME-SCALE FAULT TOLERANCE

The lazy shadowing computational model provides the basis for the design of efficient energy- and power-aware fault-tolerance approaches, that takes the energy-delay tradeoffs of the underlying



(a) Main process successful completion.



(b) Main process failure.

Figure 2: Lazy Shadowing execution dynamics.

application into consideration. In this section, we show how lazy shadowing can be used as the basic building block for the design of a resilience model for extreme-scale computing environments. The resulting model, referred to as *Leaping Shadows*, takes into consideration the main characteristics of tightly-coupled, highly-scalable applications to achieve high-tolerance to failure, while reducing energy consumption, in extreme-scale and failure-prone computing environments. In the following, we first introduce the execution and communication model of the targeted applications. We then describe the dynamics of the *Leaping Shadows* resilience model. Lastly, we discuss important aspects of its implementation.

4.1 System model

We consider the class of compute-intensive, strongly-scaled and tightly-coupled applications, executing on a large-scale multi-core computing infrastructure [2]. Communication between cores is achieved using low-latency, high-bandwidth interconnect networks, such as Infiniband.

We use W to denote the size of an application workload, and assume that the workload is split arbitrarily into a set of tasks, τ , which execute in parallel and are synchronized using barriers. Given the prominence of MPI in HPC environments, we assume message passing as the communication model between tasks. Based on this model, each pair of communicating tasks is associated with a logical first-in-first-out (FIFO) channel, which guarantees ordered delivery of messages.

Assuming that each task is assigned to execute on a core at a maximum speed $\sigma = 1$, the failure-free completion time of the application is $w = W/|\tau|$. When a failure occurs, the non-failing tasks continue executing until they reach their synchronization barrier. These tasks remain idle until the failure recovery is complete. The main objective of the *Leaping Shadows* resilience model is to minimize the idle time induced by failures and ensure forward progress

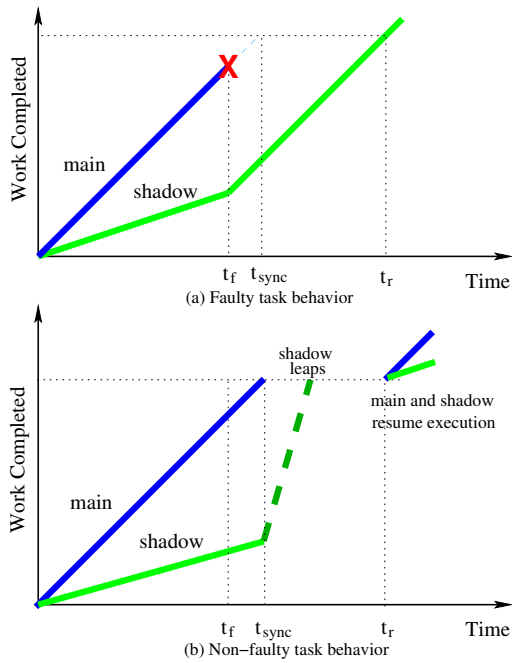


Figure 3: The concept of Leaping Shadows.

by rolling-forward shadows during the recovery phase.

4.2 Leaping Shadows

The execution of an application can be carried out by simultaneously running all main and shadow processes for all the tasks. Let $M(i)$ denote the main process executing the i^{th} task $T(i)$ of the application, and $S(i)$ its associated shadow. The task execution is divided into a set of execution phases, each extending over a synchronization interval. In the absence of failure, the behavior of a main and its lazy shadow is identical to the behavior of a main and its lazy shadow, depicted in Figure 2. Figure 3 shows the behavior of the main processes and their associated leaping shadows, assuming a failure occurrence at time t_f .

Figure 3(a) depicts the execution dynamics of a failing main and its associated shadow. Upon failure of the main process, the associated shadow continues execution, but at a higher rate, until it reaches its synchronization barrier at time t_r . Figure 3(b) illustrates the behavior of the remaining main processes and their associated shadows. Unaware of the failed process, the remaining main processes continue executing until they reach their synchronization barriers at time t_{sync} . While the shadow of the failed main process progresses toward its synchronization barrier, all remaining main processes become idle until time t_r . It is worth noting that, given the tightly-coupled and strongly-scaled nature of the application, synchronization points occur frequently. Consequently, the time between synchronization points is very small relative to the total execution time of the application. Therefore, if one main, $M(i)$, fails at time t_f , the remaining main processes will reach their synchronization point, shortly after the failure, specifically at time $t_{sync} = t_f + \epsilon$, where $\epsilon \approx 0$. The Leaping Shadows model opportunistically take advantage of the idle time, induced by the failure recovery process, to *leap forward* all shadows to their associated non-failing main processes, thereby saving execution time and energy, and guaranteeing forward progress. After the shadow of the failing main reaches its synchronization barrier, all processes re-

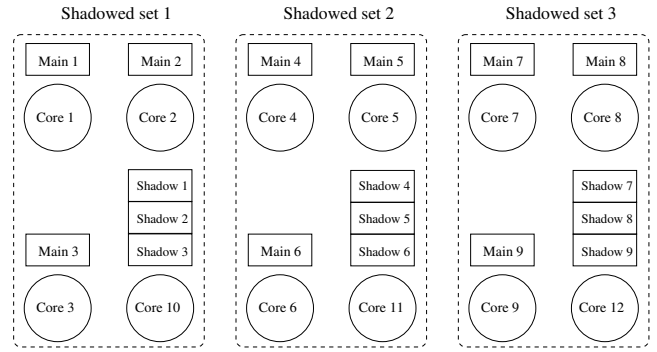


Figure 4: An example of collocation. $N = 12$, $M = 9$, $S = 3$, and $\alpha = 3$.

sume execution from a consistent synchronization point. This process continues until the completion of all tasks.

Collocation is used to control the leaping shadows execution rates. To execute an application with M tasks, $M + S$ cores are required, where M is a multiple of S , all executing at the maximum rate. M of these cores are allocated to the main processes while the remaining S cores are shared among their associated shadows. Each main process is allocated one core, while $\alpha = M/S$ shadows are collocated on a single core. α is referred to as the main to shadow ratio. For example, if $M = 9$ and $S = 3$, then the 9 shadows execute on 3 cores, with every $\alpha = 3$ shadows executing on a core, as shown in Figure 4.

Collocation of α shadows on a core has an important ramification with respect to the resilience of the system. Specifically, to speed up a shadow of a failed main to the maximum rate, all other collocated shadows must be terminated. Consequently, a second failure in any of the mains of the terminated shadows cannot be tolerated. In other words, the $M + S$ cores are grouped into S sets, which we call *shadowed sets*, each containing $\alpha + 1$ cores with α mains executing on α cores (referred to as main cores) and their corresponding α shadows collocated on one core (referred to as shadow core). Each shadowed set can tolerate a failure in any of its cores, since failure of a main core would be tolerated by the shadow core and failure of the shadow core will not affect any mains. After the first failure in a shadowed set, the set is called *vulnerable* because it cannot tolerate another failure.

The basic steps describing the execution of an application, when using Leaping Shadows based on shadow collocation, are depicted in Algorithm 1. The total workload is split into M parallel tasks (line 1). To use $M + S$ cores to execute the application, the M tasks are further grouped into S shadowed sets, each with $\alpha = M/S$ cores for α main processes and 1 core for all the associated shadow processes (line 2). The execution starts by simultaneously running all the main and shadow processes (line 3). During the execution of the application, the system runs a failure monitor that triggers corresponding actions when a failure is detected (line 5 to 18). A failure may trigger different actions, depending on its type and precedence with respect to other failures. A shadowed set becomes *vulnerable* after the occurrence of the first failure. A failure occurring in a vulnerable shadowed set (SS(j)), results in an application failure. In response, the system terminates all running processes, initiates a recovery phase, either by rebooting or replacing failing cores, and restarts execution (line 8 to 10). On the other hand, failure in a non-vulnerable shadowed set can be overcome through lazy shadowing and does not translate into application failure. In this case, the shadowed set in question (SS(j)) is marked as vulnerable (line

```

input :  $W, M, S$ 
output: Application execution status
1 split  $W$  into  $M$  tasks;
2 cluster  $M$  tasks into  $S$  shadowed sets;
3 start  $M(i)$  and  $S(i)$  for each task  $T(i)$ ;
4 while execution not done do
5   if failure detected in shadowed set  $SS(j)$  then
6     if  $SS(j)$  is vulnerable then
7       notify "Application failure";
8       terminate all mains and shadows;
9       repair all failures;
10      Restart execution;
11    else
12      mark  $SS(j)$  as vulnerable;
13      if failure happened to a main  $M(k)$  then
14        promote  $S(k)$  to  $M(k)$ ;
15        perform shadow leaping;
16      end
17    end
18  end
19 end
20 output "Application completes";

```

Algorithm 1: Lazy Shadowing

12). Note, however, that a failure in a non-vulnerable shadowed set, can be caused by the failure of a main process or its associated shadow. While a shadow failure does not impact the normal execution of the main processes, and thus can be ignored, failure of the main process forces the remaining main processes to suspend execution after they reach their synchronization point. The shadow process, $S(k)$, associated with a failing process, $M(k)$, becomes the primary process of the associated task and increases its execution to the maximum rate (line 14). This, in turn, forces the termination of all other collocated processes. Simultaneously, a forward leaping process is undertaken by all remaining shadows to align their states with that of their associated mains (line 15). This process continues until the application is successfully completed.

4.3 Leaping Shadows implementation issues

State consistency is required both during normal execution and following a failure of a main process to roll-forward the state of the leaping shadows. During normal execution, shadows remain mute, in the sense that all outgoing messages from shadows are suppressed. A shadow process, however, will typically lag behind its main process during execution. Therefore, it is necessary to ensure that the shadow's state is consistent with that of its associated main, to successfully complete its associated task in case of failure. To this end, a message-logging protocol is used to ensure consistency [4]. These protocols typically use a minimum amount of meta-information to store and replicate the non-deterministic decisions in the execution of an application. These meta-data, also called determinants, are exchanged through system-level messages. To provide correct recovery after failure, a mechanism is required to guarantee that every shadow process follows the same computation and communication steps as its main process. After a main process $M(i)$ fails, $S(i)$ will take over $M(i)$'s role to recover from this failure. If there are other shadows sharing the same core with $S(i)$, they will be terminated and $S(i)$ will start consuming the messages in its receiver-side message log at a faster speed. The message logging protocol will ensure that shadow $S(i)$ reaches a consistent state with the rest of the system.

Upon failure of a main processes, all shadow processes must update their address space to "catch up" with their associated non-failing main processes. A technology, such as remote direct memory access (RDMA), to roll-forward the state of the shadow to be consistent with that of its associated main. Rather than copying data to the buffers of the operating system, RDMA enables the network adapter to transfer data directly from the main process to its shadow. The zero-copy networking feature of RDMA considerably reduces latency, thereby enabling fast transfer of data between the main and its shadow.

5. ANALYTICAL MODELS

Three important metrics for assessing the quality of an application's execution, when Lazy Shadowing is used, are 1) the application failure probability, which measures the level of reliability; 2) expected completion time; and 3) expected energy consumption. In the following we develop mathematical models to be used for evaluation in the next section, with the understanding that process replication is a special case of Lazy Shadowing where $\alpha = 1$.

5.1 Application failure probability

Application failure, which forces the execution to start over, is inevitable even when every process is replicated. Lazy Shadowing is able to tolerate one failure in each shadowed set, and the second failure in any shadowed set implies the need to restarting the execution from the very beginning. However, Lazy Shadowing is orthogonal to checkpointing in the sense that we can combine the two, to avoid rolling the execution back to the very beginning when application failure occurs.

Since each process is replicated with a shadow, Lazy Shadowing has the potential to significantly increase the Mean Number of Failures To Interrupt (MNFTI), i.e., the average number of core failures until application failure occurs, and Mean Time To Interrupt (MTTI), i.e., the average time elapsed until application failure occurs. Therefore, the checkpointing interval should be increased to a large extent when checkpointing is combined with Lazy Shadowing. Furthermore, if the resulted checkpointing interval is larger than the completion time of the application, then checkpointing may not be used at all. In the following, the first question to study is the new MNFTI when Lazy Shadowing is used.

The impact of process replication on MNFTI has been studied in [5]. Our problem is equivalent with the difference that our work can tolerate one failure in each shadowed set while [5] can tolerate one failure in each replica-group, when each process is replicated once. Therefore, we can directly apply the methodology in [5] to our case, and the MNFTI with Lazy Shadowing for different number of shadowed sets (S) is shown in Table 1. Note that when processes are not replicated, every failure would interrupt the application, i.e., $MNFTI=1$, so MNFTI can be significantly increased by Lazy Shadowing. At the same time, it is projected that an extreme-scale application's MTTI can be increased to tens of hours from minutes assuming each core's MTBF is 25 years.

Table 1: Application's MNFTI when Lazy Shadowing is used. Results are independent of α .

S	2^0	2^1	2^2	2^3	2^4	2^5	2^6
MNFTI	3.0	3.7	4.7	6.1	8.1	11.1	15.2
S	2^7	2^8	2^9	2^{10}	2^{11}	2^{12}	2^{13}
MNFTI	21.1	29.4	41.1	57.7	81.2	114.4	161.4
S	2^{14}	2^{15}	2^{16}	2^{17}	2^{18}	2^{19}	2^{20}
MNFTI	227.9	321.8	454.7	642.7	908.5	1284.4	1816.0




Figure 5: Illustration of application's progress with failure incurred delays.

Even though the above results imply that checkpointing may not be necessary when Lazy Shadowing is used, it is important to quantify the probability that an application failure would occur during the application's execution, defined as "application failure probability". Let $f(t)$ denotes the failure probability density function of each core, and $F(t)$ be the corresponding cumulative distribution function, i.e., $F(t) = \int_0^t f(\tau) d\tau$ is the probability that a core fails in the next t time. Since each shadowed set can tolerate one failure, then the probability that a shadowed set with α main cores and 1 shadow core does not fail by time t is the probability of no failure plus the probability of one failure, i.e.,

$$P_g = (1 - F(t))^{\alpha+1} + \binom{\alpha+1}{1} F(t) \times (1 - F(t))^\alpha \quad (1)$$

and the probability that the application using N cores fails within t time is the complement of the probability that none of the shadowed sets fails, i.e.,

$$P_a = 1 - (P_g)^S \quad (2)$$

where $S = \frac{N}{\alpha+1}$ is the number of shadowed sets.

Since the application only fails during its execution, we can calculate the application failure probability using t equal to the expected completion time of the application. We will develop the model for the expected completion time in the next subsection.

5.2 Expected completion time

One of the major performance metrics of interest to end-users, is the application's completion time. To evaluate this, we develop an analytical model for the expected completion time of Lazy Shadowing, with all probabilities of failures considered. We assume that failures don't happen at the same time.

First we discuss the case of k failures, which separate the execution into $k+1$ intervals. Denote by Δ_i ($1 \leq i \leq k+1$) the i^{th} continuous execution interval, and τ_i ($1 \leq i \leq k$) the recovery time after Δ_i . The application's progress with delay incurred by failures is illustrated in Figure 5.

Since Lazy Shadowing use M cores for executing main processes and S cores for shadowing ($M + S = N$), the total workload W will be split into M tasks, each of which will be assigned a pair of main and shadow processes. Therefore, the workload of each process is $w = W/M$. The recovery time τ_i is the time needed for the lazy shadow of the failed main to catch up. With shadow leaping, it is guaranteed that all the shadows reach the same execution point as the mains (See Figure 3) after the previous recovery, so every recovery time is proportional to its previous continuous execution

length, which is Δ_i . That is, $\tau_i = \Delta_i \times (1 - \sigma_s^b)$. The value of Δ_i can be obtained given a failure probability distribution, as will be demonstrated in Section 6. Since we assume there are k failures, then Δ_{k+1} is the failure free execution interval until W is complete, i.e., $\Delta_{k+1} = w - \sum_{i=1}^k \Delta_i$. Finally, according to Figure 5, the completion time with k failures is

$$T_c^k = \sum_{i=1}^{k+1} \Delta_i + \sum_{i=1}^k \tau_i = w + (1 - \sigma_s^b) \sum_{i=1}^k \Delta_i \quad (3)$$

We assume that failures do not occur during recovery, so the failure probability of a core during the execution can be estimated as $P_c = F(w)$. Then the probability that there are k failures among the N cores during the execution is

$$P_s^k = \binom{N}{k} P_c^k (1 - P_c)^{N-k} \quad (4)$$

The completion time considering all possible failures can be averaged as $T_c = \sum_i T_c^i \cdot P_s^i$, and the expected completion time considering the possibility of re-execution after application failure is

$$T_{total} = T_c / (1 - P_a) \quad (5)$$

In the special case of $\alpha = 1$, which is process replication, half of the hardware resources are dedicated to replicas so that the workload assigned to each task is significantly increased given the fixed number of cores available, i.e., $w = 2W/N$. Different from the case of $\alpha \geq 2$, failures do not incur any delay unless application failure occurs, since the replicas are executing at the same rate as the main processes. Therefore, the completion time of process replication without application failure is constant with respect to the number of failures, and the expected completion time considering the possibility of re-execution is

$$T_{total} = T_c / (1 - P_a) = w / (1 - P_a) \quad (6)$$

A closer look at the above analysis one can realize that Lazy Shadowing has both advantage and disadvantage compared to traditional process replication. When collocating multiple shadow processes on each core, more resources will be dedicated to main processes, leading to less workload per process and thus less completion time. On the other hand, however, collocation slows down the shadow processes, implying delays when failures occur. Although it may seem that the delay would keep deteriorating as the number of failures increases, it turns out to be well bounded, as a result of shadow leaping. From Equation 3 we can see the delay of k failures is $(1 - \sigma_s^b) \sum_{i=1}^k \Delta_i$. Since we have the relation $\Delta_{k+1} = w - \sum_{i=1}^k \Delta_i$, $\sum_{i=1}^k \Delta_i$ is bounded by w , effectively limiting the delay by $(1 - \sigma_s^b)w$.

5.3 Expected energy consumption

The power consumption of one core consists of two parts, dynamic power, p_d , which exists only when the core is executing, and static power, p_s , which is constant as long as the machine is on. This can be modeled as $p = p_d + p_s$. Note that in addition to CPU leakage, other components, such as memory and disk, also contribute to the static power consumption.

For process replication, all cores are running all the time until the application is complete. Therefore, the expected energy consumption, En , is proportional to the expected execution time T_{total} :

$$En = N * p * T_{total} \quad (7)$$

Although the failed components should not consume any power, we ignore this since the number of failures is negligible compared to the total number of cores.

Lazy Shadowing has the potential to save power compared to process replication, since main cores are idle during the recovery time after each failure, and the shadows can achieve forward progress through shadow leaping. During the normal execution time, all the cores are consuming static power as well as dynamic power. During recovery time, however, the main cores are idle and consume only static power, while the shadow cores perform shadow leaping, which may lead to higher dynamic power due to memory access and communication, and then become idle with no dynamic power consumption. Again, we include the power consumption of the failed components. Altogether, the expected energy consumption for Lazy Shadowing can be modeled as

$$En = N * p_s * T_{total} + N * p_d * w + S * p_l * T_l. \quad (8)$$

with p_l denoting the dynamic power consumption during shadow leaping of each core and T_l denoting the expected total time spent on leaping during the execution of the application.

6. EVALUATION

Careful analysis of the mathematical models above leads us to identify several important factors that influence the quality of an application's execution using Lazy Shadowing. These factors can be classified into three categories, i.e., system category, application category, and Lazy Shadowing category. The system category includes static power ratio, ρ ($\rho = p_s/p$), and MTBF of each core; the application category is mainly the total workload, W ; and Lazy Shadowing category involves the number of cores to use, N , and main to shadow ratio, α , which together determine the number of main cores and shadow cores ($N = M + S$ and $\alpha = M/S$). In this section, we evaluate each performance metric of Lazy Shadowing, with the influence of each of the factors considered.

6.1 Comparison to checkpointing and process replication

To study the efficiency of Lazy Shadowing in terms of reliability, completion time, and energy consumption, we compare with traditional process replication using the analytical models developed in Section 5. We also compare to checkpointing, of which the completion time is calculated with Daly's model [9] assuming the recovery time and checkpointing time are both 10 minutes, and then the energy consumption is derived using Equation 7. It is clear from Equation 3 that the total recovery delay $\sum_{i=1}^k \tau_i$ is determined by the execution time $\sum_{i=1}^k \Delta_i$, independent of the distribution of failures which determines the individual value of Δ_i . Therefore, our models are generic with no assumption about failure probability distribution, and the expectation of the total delay from all failures is the same as the failures are uniformly distributed [9]. Specifically, $\Delta_i = w/(k+1)$, and T_c^k can be rewritten as $w + w * (1 - \sigma_s^b) * \frac{k}{k+1}$ for Lazy Shadowing. Further, we assume that each shadow process gets a fair share of its shadow core's execution rate so that $\sigma_s^b = \frac{1}{\alpha}$. One may argue that the execution rate of the shadows should be degraded because of increased memory pressure or communication overhead. To take that into account, we have also studied the slowing down effect using a penalty model, to be discussed later. To calculate the expected energy consumption for Lazy Shadowing with Equation 8, we further assume that the dynamic power during shadow leaping is twice that during normal execution, i.e., $p_l = 2 * p_d$, and the time for shadow leaping through RDMA is half of the time required for recovery, i.e., $T_l = 0.5 * (T_{total} - w)$. The static power ratio ρ is fixed at 0.5 for now, other values will be explored later in this section.

The first study uses $N = 1$ million cores, effectively simulating the future extreme-scale computing environment, and assumes that

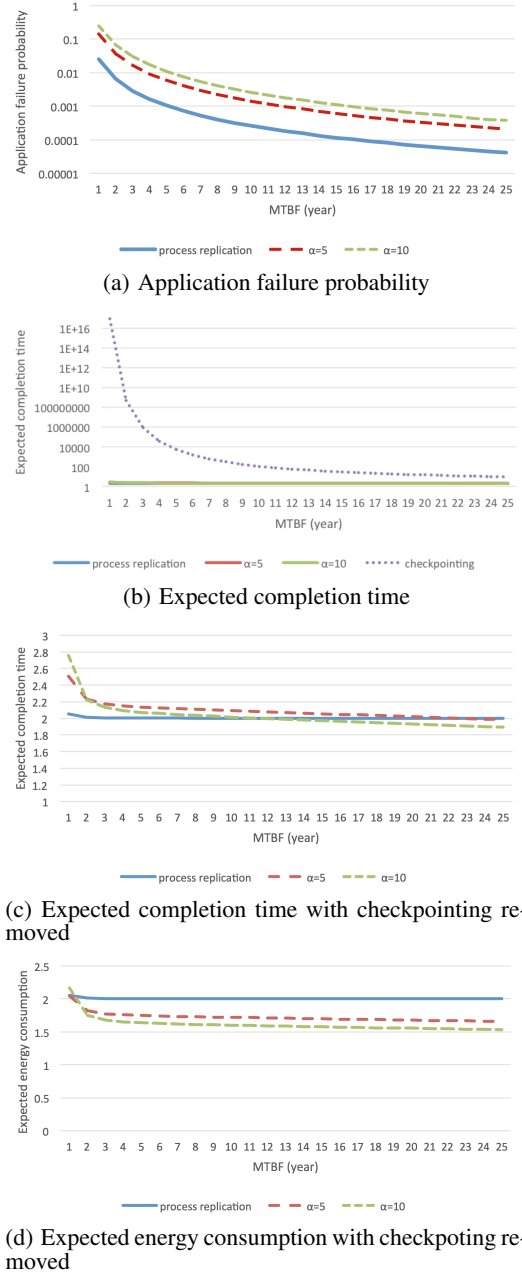


Figure 6: Comparison to process replication and checkpointing. $W = 10^6$ hours, $N = 10^6$, $\rho = 0.5$.

$W = 1$ million hours. For Lazy Shadowing we varied α from 1 to 10, with the understanding that it is unrealistic to collocate too many processes on a core. Besides $\alpha = 1$, which is equivalent to process replication, we only show $\alpha = 5$ and $\alpha = 10$ as others can be easily inferred from the figures.

By definition, the application failure probability for checkpointing is 0, as it periodically saves the execution states, from which the computation can be resumed upon failure. Figure 6(a) shows that due to collocation, the application failure probability increases for Lazy Shadowing. However, even with the lowest MTBF (1 year), Lazy Shadowing is still able to complete the application without re-execution with larger than 0.75 probability.

It is clear from Figure 6(b) that checkpointing incurs significant delay, indicating that it may not be a viable fault tolerance approach for future extreme-scale computing. The completion time and energy consumption of checkpointing are so large that the comparison between process replication and Lazy Shadowing is invisible. Therefore, we re-plot the comparison between Lazy Shadowing and process replication in Figure 6(c) and 6(d), with checkpointing removed. Figure 6(c) reveals that the most time efficient choice depends on MTBF. More specifically, process replication is more suited when MTBF is low while otherwise Lazy Shadowing is better. In terms of energy consumption, the advantage of Lazy Shadowing is much more obvious. For MTBF from 2 to 25 years, Lazy Shadowing with $\alpha = 5$ can achieve 9.6-17.1% energy saving, while the saving increases to 13.1- 23.3% for $\alpha = 10$. The only exception is when MTBF is extremely low (1 year), Lazy Shadowing with $\alpha = 10$ consumes more energy than process replication because of its higher probability of re-execution.

6.2 Impact of the number of cores

The system scale, measured in number of cores, has a direct impact on the failure rate seen by the application. To study its impact, we varies N from 10,000 to 1,000,000 with W scaled proportionally, i.e., $W = N$. When MTBF is 5 years, the results are shown in Figure 7. Please note that the time and energy for checkpointing when $N = 1,000,000$ are beyond the scope of the figures, so we mark their values on top of their columns. When completion time is considered, Figure 7(a) clearly shows that each of the three fault tolerance alternatives has its own advantage. Specifically, checkpointing is the best choice for small systems with around 10,000 cores, Lazy Shadowing outperforms others for systems with around 100,000 cores, while process replication has slight advantage over Lazy Shadowing for larger systems. On the other hand, Lazy Shadowing wins for all system sizes covered in Figure 7(b), when energy consumption is considered. The comparison is different when MTBT is 25 years, as shown in Figure 8. Though improved to a large extent, the performance of checkpointing is still much worse than that of the other two approaches when N is 1,000,000. Lazy Shadowing also benefits from the increased MTBF, and further reduces its completion time and energy consumption for system sizes. In addition, Lazy Shadowing is able to achieve shorter completion time than process replication when N reaches 1,000,000.

6.3 Impact of workload

To a large extent, workload determines the time exposed to failure. With other factors being the same, an application with a larger workload is likely to encounter more failures during its execution. Hence, it is intuitive that workload would impact the performance comparison among the three fault tolerance approaches.

Fixing N at 1,000,000, we increase W from 1,000,000 hours to 12,000,000 hours. Figure 9 assumes a MTBF of 25 years and shows both the time and energy. Checkpointing has the worst performance in both time and energy in all cases. In terms of completion time, process replication is more efficient than Lazy Shadowing when workload reaches 6,000,000 hours. Considering energy consumption, however, Lazy Shadowing is able to achieve the most energy saving in all cases. When MTBF of 5 years is used, the difference is that process replication consumes less energy than Lazy Shadowing when W reaches 6,000,000.

6.4 Impact of static power ratio

With various architectures and organizations, servers deployed at different supercomputers have different characteristics in terms of power consumption. The static power ratio ρ is used to abstract

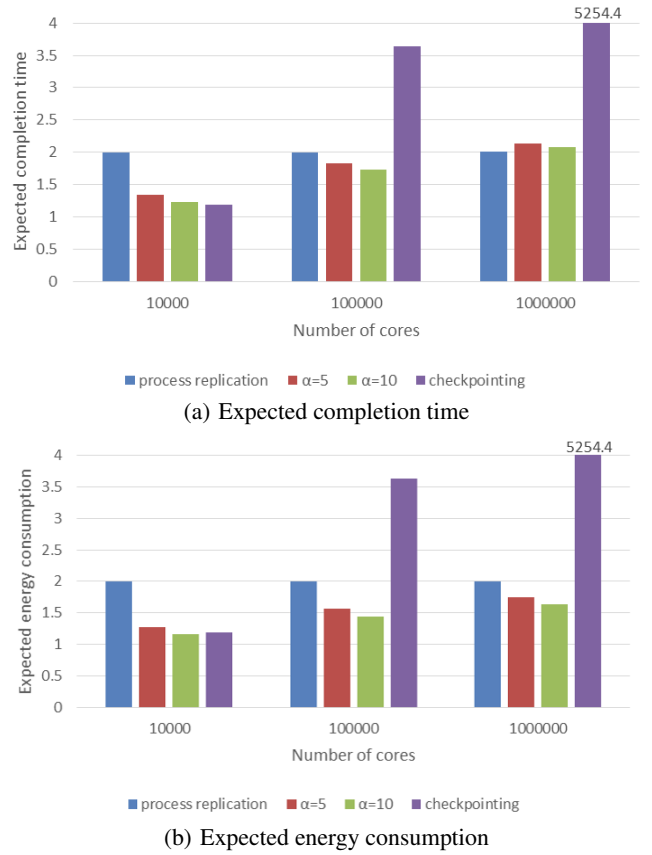



Figure 7: Sensitivity to number of cores. $W = N$, MTBF=5 years, $\rho = 0.5$.


the amount of static power consumed versus dynamic power. In our models, ρ does not impact the completion time, but power and energy consumption. Considering modern systems, we vary ρ from 0.3 to 0.7 and study its effect on the expected energy consumption. The results for Lazy Shadowing with $\alpha = 5$ are normalized to that of process replication and shown in Figure 10. The results for other values of α have similar behavior and thus are not shown. Lazy Shadowing achieves more energy saving when static power ratio is low, since it saves dynamic power but not static power. When static power ratio is low ($\rho = 0.3$), Lazy Shadowing is able to save 20%-24% energy for the MTBF of 5 to 25 years. The saving decreases to 5%-11% when ρ reaches 0.7.

6.5 Adding collocation overhead

The last study is conducted to capture the impact on the performance of Lazy Shadowing brought by collocation overhead. We re-model the speed of shadows as $\sigma_s^b = \frac{1}{\alpha^{1.5}}$ to simulate the effect of memory thrashing and context switch. Figure 11 shows the impact of collocation overhead on expected energy consumption for Lazy Shadowing with $\alpha = 5$, with all the values normalized to that of process replication. The results for other values of α have similar behavior and thus are not shown. As expected, energy consumption is penalized because of slowing down of the shadows. It is surprising, however, that the impact is negligible, with the largest difference being 4.4%. The reason is that Lazy Shadowing can take advantage of the recovery time after each failure and achieve forward progress for shadow processes that fall behind. When $\alpha = 10$, the largest difference further decreases to 2.5%.



(a) Expected completion time




(b) Expected energy consumption

Figure 8: Sensitivity to number of cores. $W = N$, MTBF=25years, $\rho = 0.5$.


7. CONCLUSION AND FUTURE WORK

As the scale and complexity of HPC continue to increase, both the failure rate and energy consumption are expected to increase dramatically, making it extremely challenging to deliver extreme-scale computing that is efficient and reliable. Existing fault tolerance methods rely on either time or hardware redundancy. Neither of them appeals to the next generation of supercomputing, as the first choice may incur significant delay while the second one constantly wastes over 50% of the system efficiency. The need for an efficient and reliable solution in extreme-scale, failure-prone computing environments calls for a reconsideration of the fault tolerance problem.

To this end, we propose a proactive, power-aware resilience model, referred to as Lazy Shadowing, as an efficient and scalable alternative to achieve high-levels of fault tolerance for future extreme-scale computing. In this paper, we present a comprehensive discussion of the techniques that will enable Lazy Shadowing. In addition, we develop a series of mathematical models to assess its performance in terms of reliability, completion time, and energy consumption, with all the critical parameters identified. Through comparison with existing fault tolerance approaches, we identify the scenarios where each of the alternatives should be chosen. Specially, checkpointing consumes the least time and energy when system size is small; process replication should be used to minimize completion time when system size is extremely large and failure rate is extremely high; and Lazy Shadowing is the choice for all other cases, for both completion time and energy consumption. In



(a) Expected completion time



(b) Expected energy consumption

Figure 9: Sensitivity to workload. $N = 10^6$, MTBF=25 years, $\rho = 0.5$.




Figure 10: Impact of static power ratio on energy consumption normalized to process replication. $W = 10^6$ hours, $N = 10^6$, $\alpha=5$.

addition, we predict that Lazy Shadowing is able to achieve over 20% of energy saving with potential completion time reduction for future extreme-scale computing. In the future, we plan to implement Lazy Shadowing and measure its performance in reality.

8. REFERENCES

- [1] S. Agarwal, R. Garg, M. S. Gupta, and J. E. Moreira. Adaptive incremental checkpointing for massively parallel systems. In *ICS*, St. Malo, France, 2004.

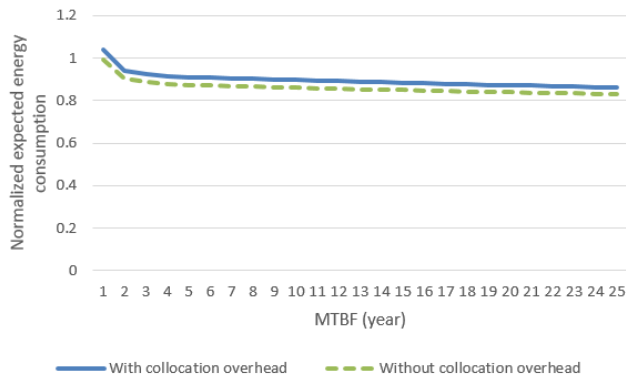


Figure 11: Impact of collocation overhead on normalized energy consumption. $W = 10^6$ hours, $N = 10^6$, $\rho=0.5$, $\alpha=5$.

- [2] S. Ahern and et. al. Scientific discovery at the exascale, a report from the doe ascr 2011 workshop on exascale data management, analysis, and visualization, 2011.
- [3] L. Alvisi and et. al. An analysis of communication induced checkpointing. In *Fault-Tolerant Computing*, 29th International Symposium on, pages 242–249, 1999.
- [4] L. Alvisi and K. Marzullp. Message logging: Pessimistic, optimistic, causal, and optimal. *IEEE Trans. Softw. Eng.*, 42(2):149–159, 1998.
- [5] H. Casanova, Y. Robert, F. Vivien, and D. Zaidouni. Combining Process Replication and Checkpointing for Resilience on Exascale Systems. Rapport de recherche RR-7951, INRIA, May 2012.
- [6] K. Chandy and C. Ramamoorthy. Rollback and recovery strategies for computer programs. *Computers, IEEE Transactions on*, C-21(6):546–556, June 1972.
- [7] H. chang Nam, J. Kim, S. Lee, and S. Lee. Probabilistic checkpointing. In *In Proceedings of Intl. Symposium on Fault-Tolerant Computing*, pages 153–160, 1997.
- [8] F. Chen, D. A. Koufaty, and X. Zhang. Hystor: Making the best use of solid state drives in high performance storage systems. ICS, pages 22–32, New York, USA, 2011.
- [9] J. Daly. A higher order estimate of the optimum checkpoint interval for restart dumps. *Future Gener. Comput. Syst.*, 22(3):303–312, Feb. 2006.
- [10] J. Elliott and et. al. Combining partial redundancy and checkpointing for HPC. ICDCS, pages 615–626, Washington, DC, USA, 2012. IEEE Computer Society.
- [11] E. Elnozahy and et. al. A survey of rollback-recovery protocols in message-passing systems. *ACM Comput. Surv.*, 34(3):375–408, 2002.
- [12] E. Elnozahy and J. Plank. Checkpointing for peta-scale systems: a look into the future of practical rollback-recovery. *Dependable and Secure Computing, IEEE Transactions on*, 1(2):97–108, april-june 2004.
- [13] E. Elnozahy and W. Zwaenepoel. Manetho: Transparent rollback-recovery with low overhead, limited rollback and fast output commit. *TC*, 41:526–531, 1992.
- [14] D. Fiala and et. al. Detection and correction of silent data corruption for large-scale high-performance computing. SC, pages 78:1–78:12, Los Alamitos, CA, USA, 2012.
- [15] A. Guermouche and et. al. Uncoordinated checkpointing without domino effect for send-deterministic mpi applications. In *IPDPS*, pages 989–1000, May 2011.
- [16] D. Hakkarinen and Z. Chen. Multilevel diskless checkpointing. *Computers, IEEE Transactions on*, 62(4):772–783, April 2013.
- [17] P. Hargrove and J. Duell. Berkeley lab checkpoint/restart (blcr) for linux clusters. In *Journal of Physics: Conference Series*, volume 46, page 494, 2006.
- [18] J. Helary and et. al. Preventing useless checkpoints in distributed computations. In *Reliable Distributed Systems*, pages 183–190, oct 1997.
- [19] S. Kalaiselvi and V. Rajaraman. A survey of checkpointing algorithms for parallel and distributed computers. *Sadhana*, 25(5):489–510, 2000.
- [20] L. Lamport. The part-time parliament. *ACM Trans. Comput. Syst.*, 16(2):133–169, May 1998.
- [21] K. Li, J. F. Naughton, and J. S. Plank. Low-latency, concurrent checkpointing for parallel programs. *IEEE Trans. Parallel Distrib. Syst.*, 5(8):874–879, Aug. 1994.
- [22] B. Mills and et. al. Shadow computing: An energy-aware fault tolerant computing model. In *ICNC*, 2014.
- [23] A. Moody, G. Bronevetsky, K. Mohror, and B. Supinski. Design, modeling, and evaluation of a scalable multi-level checkpointing system. In SC, pages 1–11, 2010.
- [24] X. Ni, E. Meneses, N. Jain, and L. V. Kalé. Acr: Automatic checkpoint/restart for soft and hard error protection. SC, pages 7:1–7:12, New York, NY, USA, 2013. ACM.
- [25] J. Plank and K. Li. Faster checkpointing with n+1 parity. In *Fault-Tolerant Computing*, pages 288–297, June 1994.
- [26] J. Plank and M. Thomason. The average availability of parallel checkpointing systems and its importance in selecting runtime parameters. In *Fault-Tolerant Computing*, pages 250–257, 1999.
- [27] B. Randell. System structure for software fault tolerance. In *Proceedings of the international conference on Reliable software*, pages 437–449, New York, NY, USA, 1975. ACM.
- [28] R. Riesen, K. Ferreira, J. R. Stearley, R. Oldfield, J. H. L. III, K. T. Pedretti, and R. Brightwell. Redundant computing for exascale systems, December 2010.
- [29] F. B. Schneider. Implementing fault-tolerant services using the state machine approach: A tutorial. *ACM Comput. Surv.*, 22(4):299–319, Dec. 1990.
- [30] J. Stearley and et. al. Does partial replication pay off? In *DSN-W*, pages 1–6, June 2012.
- [31] G. Zheng, L. Shi, and L. Kale. FTC-Charm++: an in-memory checkpoint-based fault tolerant runtime for Charm++ and MPI. In *Cluster Computing, 2004 IEEE International Conference on*, pages 93–103, Sept 2004.